

# 1000-719bMSB Modeling of Complex Biological Systems

## Deep Neural Network: Unsupervised Learning

### Variational Autoencoder on MNIST

VAEs are directed probabilistic graphical models (DPGM) whose posterior is approximated by a neural network, forming an autoencoder-like architecture. Unlike discriminative modeling that aims to learn a predictor given observation, generative modeling tries to learn how the data is generated, and to reflect the underlying causal relations. [...] Variational autoencoder models make strong assumptions concerning the distribution of latent variables. They use a variational approach for latent representation learning, which results in an additional loss component and a specific estimator for the training algorithm called the Stochastic Gradient Variational Bayes (SGVB) estimator. [Wikipedia](#)

We build a VAE and train it on the MNIST dataset.

```
In [2]: ##tensorflow version 1.x
import tensorflow as tf
print(tf.__version__)
tf.compat.V1.disable_eager_execution()
import keras
from keras import layers
from keras import backend as K
from keras.models import Model
from keras.datasets import mnist
import matplotlib.pyplot as plt
import numpy as np

img_shape = (28, 28, 1)
batch_size = 16
latent_dim = 2

2.7.0
```

```
In [3]: #Define encoder network
#Note that we are using keras functional API
input_img = keras.Input(shape=img_shape)
x = layers.Conv2D(32, 3, padding='same', activation='relu')(input_img)
x = layers.Conv2D(64, 3, padding='same', activation='relu', strides=(2, 2))(x)
x = layers.Conv2D(64, 3, padding='same', activation='relu')(x)
x = layers.Conv2D(64, 3, padding='same', activation='relu')(x)

shape_before_flattening = K.int_shape(x)
x = layers.Flatten()(x)
x = layers.Dense(32, activations='relu')(x)
#x = layers.Dense(2, activation='relu')(x)

z_mean = layers.Dense(latent_dim)(x)
z_log_var = layers.Dense(latent_dim)(x)
```

```
In [4]: #Sampling from the distributions to obtain latent space
def sampling(args):
    z_mean, z_log_var = args
    epsilon = K.random_normal(shape=(K.shape(z_mean)[0], latent_dim),
                               mean=0., stddev=1.)
    return z_mean + K.exp(z_log_var) * epsilon
z = layers.Lambda(sampling)([z_mean, z_log_var])

encoder = Model(input_img, z)
```

```
In [5]: #Define decoder network
decoder_input = layers.Input(K.int_shape(z)[1:])
x = layers.Dense(np.prod(shape_before_flattening[1:]), activation='relu')(decoder_input)
x = layers.Reshape(shape_before_flattening[1:])(x)
x = layers.Conv2DTranspose(32, 3, padding='same', activation='relu', strides=(2, 2))(x)
x = layers.Conv2D(1, 3, padding='same', activation='sigmoid')(x)

decoder = Model(decoder_input, x)
z_decoded = decoder(z)
```

```
In [6]: def vae_loss(input_img, z_decoded):
    input_img = K.flatten(input_img)
    z_decoded = K.flatten(z_decoded)
    xent_loss = keras.metrics.binary_crossentropy(input_img, z_decoded)
    kl_loss = -5e-4 * K.mean(1 + z_log_var - K.square(z_mean) - K.exp(z_log_var), axis=-1)
    return K.mean(xent_loss + kl_loss)
```

```
In [7]: vae = Model(input_img, z_decoded)
vae.compile(optimizer='adam', loss=vae_loss)
```

```
In [8]: vae.summary()
decoder.summary()
```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 28, 28, 1)	0	[]
conv2d (Conv2D)	(None, 28, 28, 32)	320	['input_1[0][0]']
conv2d_1 (Conv2D)	(None, 14, 14, 64)	18496	['conv2d[0][0]']
conv2d_2 (Conv2D)	(None, 14, 14, 64)	36928	['conv2d_1[0][0]']
conv2d_3 (Conv2D)	(None, 14, 14, 64)	36928	['conv2d_2[0][0]']
Flatten (Flatten)	(None, 12544)	0	['conv2d_3[0][0]']
dense (Dense)	(None, 32)	401440	['flatten[0][0]']
dense_1 (Dense)	(None, 2)	66	['dense[0][0]']
dense_2 (Dense)	(None, 2)	66	['dense[0][0]']
lambda (Lambda)	(None, 2)	0	['dense_1[0][0]', 'dense_2[0][0]']
model_1 (Functional)	(None, 28, 28, 1)	56385	['lambda[0][0]']

-----  
Total params: 550,629  
Trainable params: 550,629  
Non-trainable params: 0

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 2)	0
dense_3 (Dense)	(None, 12544)	37632
reshape (Reshape)	(None, 14, 14, 64)	0
conv2d_transpose (Conv2DTranspose)	(None, 28, 28, 32)	18464
conv2d_4 (Conv2D)	(None, 28, 28, 1)	289

-----  
Total params: 56,385  
Trainable params: 56,385  
Non-trainable params: 0

```
In [9]: #load the data and split into train + test sets
(x_train, _), (x_test, y_test) = mnist.load_data()
```

```
In [10]: x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.astype('float32') / 255.
x_test = x_test.reshape(x_test.shape + (1,))
```

We are ready to train VAE on this dataset

```
In [11]: vae.fit(x=x_train,y=x_train, shuffle=True, epochs=10, batch_size=batch_size)
```

```
Train on 60000 samples
Epoch 1/10
60000/60000 [=====] - 182s 3ms/sample - loss: 0.2272
Epoch 2/10
60000/60000 [=====] - 174s 3ms/sample - loss: 0.2013
Epoch 3/10
60000/60000 [=====] - 177s 3ms/sample - loss: 0.1947
Epoch 4/10
60000/60000 [=====] - 197s 3ms/sample - loss: 0.1909
Epoch 5/10
60000/60000 [=====] - 208s 3ms/sample - loss: 0.1885
Epoch 6/10
60000/60000 [=====] - 183s 3ms/sample - loss: 0.1865
Epoch 7/10
60000/60000 [=====] - 185s 3ms/sample - loss: 0.1851
Epoch 8/10
60000/60000 [=====] - 182s 3ms/sample - loss: 0.1840
Epoch 9/10
60000/60000 [=====] - 178s 3ms/sample - loss: 0.1831
Epoch 10/10
60000/60000 [=====] - 174s 3ms/sample - loss: 0.1824
Out[11]: <keras.callbacks.History at 0x1de7140610>
```

```
In [12]: # to save the trained VAE on Google Drive and load them later. Not necessary
from google.colab import drive
drive.mount('/content/drive')
```

```
-----
ModuleNotFoundError Traceback (most recent call last)
<ipython-input-12-de012179945> in <module>
----> 1 # to save the trained VAE on Google Drive and load them later. Not necessary
-----> 2 from google.colab import drive
      3 drive.mount('/content/drive')
ModuleNotFoundError: No module named 'google.colab'
```

```
In [ ]: vae.save_weights('drive/My Drive/Colab Notebooks/vae_weights.h5')
vae.load_weights('drive/My Drive/Colab Notebooks/vae_weights.h5')
```

```
In [13]: # x_test is data
# encoded: means we are mapping our data point (x_test) into latent space; give us a point in latent space
encoded = encoder.predict(x_test)
```

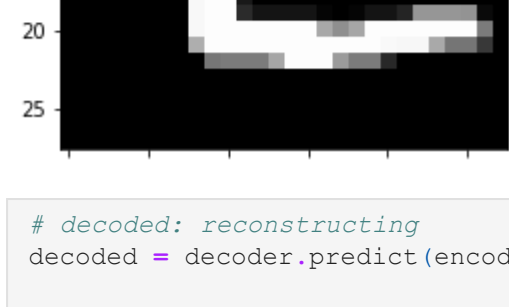
```
C:\Users\zozsia\anaconda3\lib\site-packages\keras\engine\training_v1.py:2079: UserWarning: 'Model.state_updates' will be removed in a future version. This property should not be used in TensorFlow 2.0, as 'updates' are applied automatically.
  updates=state_updates,
```

```
In [14]: print(encoded)
```

```
[ [ 2.36461 -1.4998004 ]
 [ 0.04451078 -0.02951086 ]
 [-1.3342675 -3.2245126 ]
 ...
 [ 1.2987785 -0.90321654 ]
 [-0.9657645 -0.82044506 ]
 [-0.5781946 0.19804764 ] ]
```


```
In [15]: print(encoded[0])
plt.matshow(x_test[0][:,:], cmap='gray')
```

```
Out[15]: <matplotlib.image.AxesImage at 0x1de887b910>
```



```
In [16]: print(encoded[1])
plt.matshow(x_test[1][:,:], cmap='gray')
```


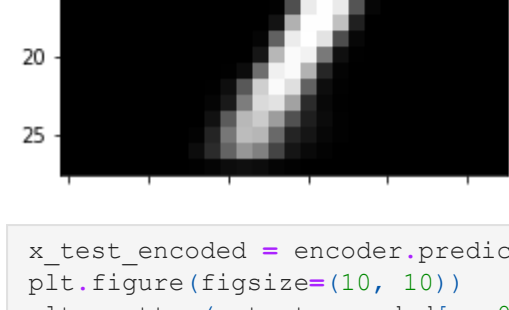
```
Out[16]: <matplotlib.image.AxesImage at 0x1de894da90>
```




```
In [17]: # decoded: reconstructing
decoded = decoder.predict(encoded)
#decoded_somc_random = decoder.predict()
```

```
In [18]: plt.matshow(x_test[0][:,:], cmap='gray')
plt.matshow(decoded[0][:,:], cmap='gray')
```

```
Out[18]: <matplotlib.image.AxesImage at 0x1de9b64400>
```



```
In [19]: x_test_encoded = encoder.predict(x_test)
plt.figure(figsize=(10, 10))
plt.scatter(x_test_encoded[:, 0], x_test_encoded[:, 1], c=y_test, cmap='jet')
plt.colorbar()
plt.show()
```



### HOMEWORK 1

Interpolate between two latent vectors (i.e., moving in the latent space) and decode and visualize the interpolations. Provide a code (python notebook) to do this interpolation; and provide 3 interpolated images in-between.

```
In [20]: def interpolate(vec1, vec2, nb):
    linspace = np.linspace(0, 3, num=nb)
    vectors = []
    for lin in linspace:
        v_new = (1.0 - lin) * vec1 + lin * vec2
        vectors.append(v_new)
    return np.asarray(Vectors)

#selecting pictures
nb1 = 2
nb2 = 0
encoded_1 = encoded[nb1]
encoded_2 = encoded[nb2]


# interpolating between chosen pictures and decoding them
interpolated_vector = interpolate(encoded_1, encoded_2, 3)
decoded_interpolation = decoder.predict(interpolated_vector)

# Printing the results
print('Original image from class ' + str(nb1))
plt.matshow(x_test[nb1][:,:], cmap='gray')
plt.show()


print('Original image from class ' + str(nb2))
plt.matshow(x_test[nb2][:,:], cmap='gray')
plt.show()

print('Interpolated images')
for index, param in enumerate(decoded_interpolation):
    plt.matshow(index)\n'.format(index+1))
    plt.matshow(decoded_interpolation[index][:,:], cmap='gray')
    plt.show()
```

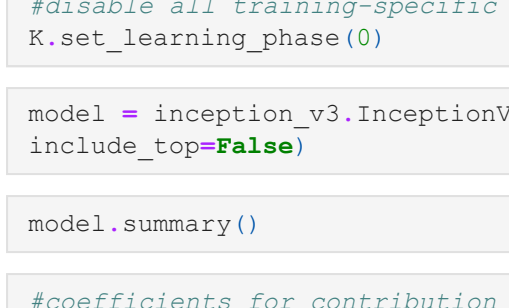
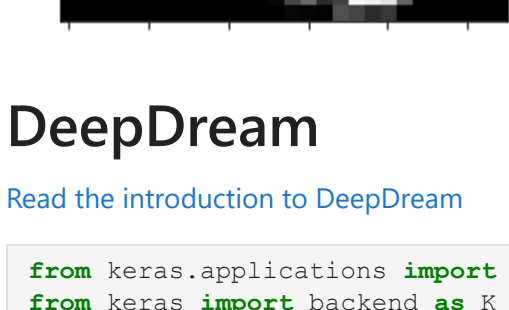



Original image from class 2



Original image from class 0



Interpolated images



### DeepDream

[Read the introduction to DeepDream](#)

```
In [ ]: from keras.applications import inception_v3
from keras import backend as K
#disable all training-specific options
K.set_learning_phase(0)
```

```
In [ ]: model = inception_v3.InceptionV3(weights='imagenet',
include_top=False)
```

```
In [ ]: model.summary()
```

```
In [ ]: #coefficients for contribution of selected layers
layer_contributions = {
    'mixed2': 0.2,
    'mixed3': 3.,
    'mixed4': 2.,
    'mixed5': 1.5,
}
```

```
In [ ]: #Creates a dictionary that maps layer names to layer instances
layer_dict = dict([(layer.name, layer) for layer in model.layers])
```

```
In [ ]: layer_dict['mixed2']
```

```
In [ ]: #define the loss
loss = K.variable(0.)
for layer_name in layer_contributions:
    coeff = layer_contributions[layer_name]
    activation = layer_dict[layer_name].output
    scaling = K.prod(K.cast(K.shape(activation), 'float32'))
    #here the loss is the L2 norm of activations of whole layer
    #before: only chosen filter
    loss = loss + coeff * K.sum(K.square(activation[:, 2:-2, 2:-2, :])) / scaling
```

```
In [ ]: #define gradient-ascent process

#this was before a noise image, now image of choice
dream = model.input

#gradient of loss wrt input image
grads = K.gradients(loss, dream)[0]
#normalize gradient
grads /= K.maximum(K.mean(K.abs(grads)), 1e-7)

outputs = [loss, grads]
fetch_loss_and_grads = K.function([dream], outputs)

def eval_loss_and_grads(x):
    outs = fetch_loss_and_grads([x])
    loss_value = outs[0]
    grad_values = outs[1]
    return loss_value, grad_values

def gradient_ascent(x, iterations, step, max_loss=None):
    for i in range(iterations):
        loss_value, grad_values = eval_loss_and_grads(x)
        if max_loss is not None and loss_value > max_loss:
            break
        x = x + step * grad_values
    return x
```

```
In [ ]: #Define some auxiliary function
import scipy
from keras.preprocessing import image
import imageio

def resize_img(img, size):
    img = np.copy(img)
    factors = (1.,)
    float_size[0] / img.shape[1],
    float_size[1] / img.shape[2],
    1)
    return scipy.ndimage.zoom(img, factors, order=1)

def save_img(img, fname):
    pil_img = deprocess_image(np.copy(img))
    imageio.imwrite(fname, pil_img)

def preprocess_image(image_path):
    img = image.load_img(image_path, target_size=(400,400))
    img = image.img_to_array(img)
    img = np.expand_dims(img, axis=0)
    img = inception_v3.preprocess_input(img)
    return img

def deprocess_image(x):
    if K.image_data_format() == 'channels_first':
        x = x.reshape((3, x.shape[2], x.shape[3]))
        x = x.transpose((1, 2, 0))
    else:
        #undoes the preprocessing done by inception_v3.preprocess_input
        x = x.reshape((x.shape[1], x.shape[2], 3))
        x /= 2.
        x *= 0.5
        x += 255.
        x = np.clip(x, 0, 255).astype('uint8')
    return x
```

```
In [ ]: base_image_path = 'drive/My Drive/Colab Notebooks/hill_wiki.jpg'
img = preprocess_image(base_image_path)
```

```
In [ ]: step = 0.01
iterations = 20
max_loss = 10.

img = gradient_ascent(img,
    iterations=iterations,
    step=step,
    max_loss=max_loss)

plt.rcParams['figure.figsize'] = (10,10)
plt.imshow(deprocess_image(img))
```

```
In [ ]: import numpy as np
step = 0.01
num_octave = 3
octave_scale = 1.4
iterations = 20
max_loss = 10.
base_image_path = 'drive/My Drive/Colab Notebooks/hill_wiki.jpg'
img = preprocess_image(base_image_path)
```

```
In [ ]: original_shape = img.shape[1:3]
successive_shapes = [original_shape]
for i in range(1, num_octave):
    shape = tuple([int(dim / (octave_scale ** i))
                   for dim in original_shape])
    successive_shapes.append(shape)

successive_shapes = successive_shapes[::-1]
original_img = np.copy(img)
shrunk_original_img = resize_img(img, successive_shapes[0])
```

```
In [ ]: successive_shapes
```

```
In [ ]: #to make the effect look cooler
for shape in successive_shapes:
    print('Processing image shape', shape)
    #upscale image and apply gradient descent
    img = resize_img(img, shape)
    img = gradient_ascent(img,
        iterations=iterations,
        step=step,
        max_loss=max_loss)

    #insert details lost by upscaling
    upscaled_shrunk_original_img = resize_img(shrunk_original_img, shape) #will be pixelated
    same_size_original = resize_img(original_img, shape) #high-quality image by downscaling or
    lost_detail = same_size_original - upscaled_shrunk_original_img #detail that was lost when scaling up
    img = img + lost_detail * alpha * insert_rhe_detail
    shrunk_original_img = resize_img(original_img, shape)
```

```
In [ ]: plt.imshow(deprocess_image(img))
```

```
In [ ]: 
```